

# Brain-Inspired Learning Machines

## Neural Models of Associative Memory and Working Memory

Emre Neftci

Department of Cognitive Sciences, UC Irvine,

November 18, 2016

## Types of Declarative Memory

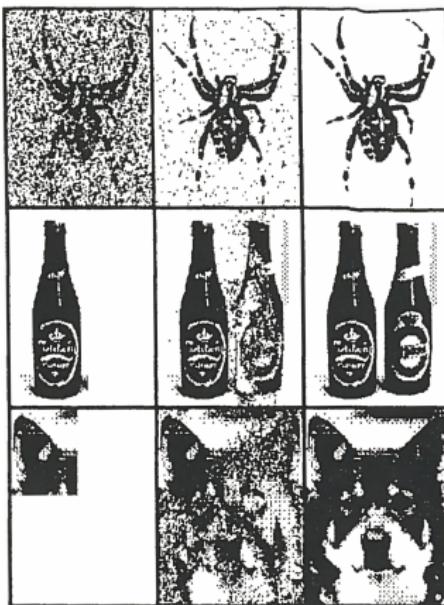
- **Working Memory** (<1m) - Prefrontal lobe, *Example:* Remembering digits of a phone number
- **Long-term Memory** (years) - Hippocampus, Inferior Temporal (IT) lobe, *Example:* Recognition Memory, Autoassociative memory

## Associative Memory

## Associative Memory

"In associative memory, a partial or approximate representation of a stored item is used to recall the full item"

Dayan and Abbott,, 2001



Hertz, Krogh, and Palmer,, 1991

In computer science, autoassociative memory is called Content-Addressable Memory (CAM)

## Address-based Memory vs. Content-based Memory

Address-based memory e.g. Random Access Memory (RAM):

- High capacity
- Need to know the “address” to retrieve memory
- Recall is not robust/fault-tolerant

Content-based memory e.g. Cache:

- Low capacity
- Ability to recall a memory for partial cues
- Errors in recall cues can be corrected
- Hardware can be fault tolerant

Biological memory mechanisms for long-term storage are CAMs

(a) Stored Memories

moscow-----russia
lima-----peru
london-----england
tokyo-----japan
edinburgh-scotland
ottawa-----canada
oslo-----norway
stockholm---sweden
paris-----france

(b) Recall

moscow---	:::::::::::   ⇒  moscow-----russia
:::::::::::canada   ⇒  ottawa-----canada	
otowa-----canada   ⇒  ottawa-----canada	
egindurrh-sxotland   ⇒  edinburgh-scotland	

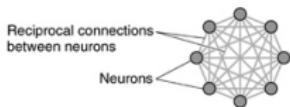
(a) Stored Memories



(b) Recall



# Hebb's Cell Assembly

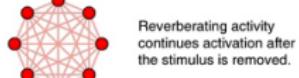


(a) The cell assembly

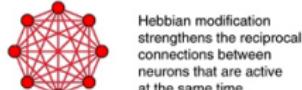


External stimulus

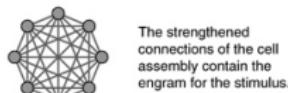
Activation of the cell assembly by a stimulus.



Reverberating activity continues activation after the stimulus is removed.



Hebbian modification strengthens the reciprocal connections between neurons that are active at the same time.



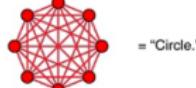
The strengthened connections of the cell assembly contain the engram for the stimulus.

(b)

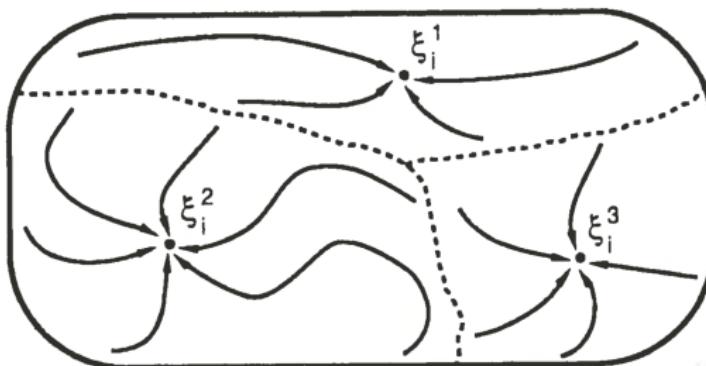


After learning, partial activation of the assembly leads to activation of the entire representation of the stimulus.

(c)



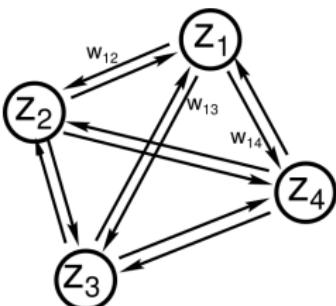
= "Circle."



- The rectangle is the **state space**.
- $\xi_i^1$ ,  $\xi_i^2$ ,  $\xi_i^3$  are **attractors**
- Regions bounded by dashed lines represent **basins of attraction**.

# Artificial Attractor Neural Network

## Hopfield Network



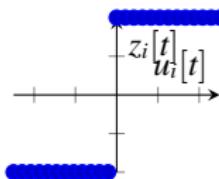
A Hopfield network is a recurrent neural net for content-addressable memory

Hopfield, *Proceedings of the national academy of sciences*, 1982

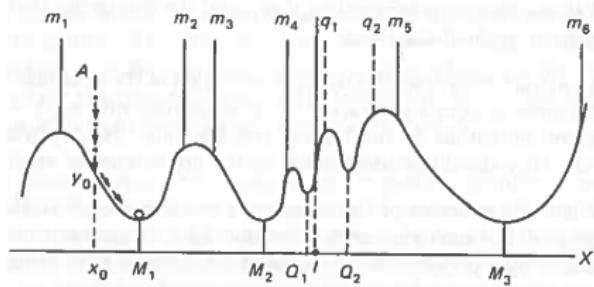
**Neuron model:**

$$u_i[t] = \sum_j w_{ij} z_j[t] + \theta_i$$

$$z_i[t] = \Theta(u_i[t]) = \begin{cases} 1 & \text{if } u_i[t] > 0 \\ -1 & \text{if } u_i[t] \leq 0 \end{cases}$$



## Energy landscape metaphor and learning



Amit., 1992

### “Energy” of a Hopfield Network:

$$E(\mathbf{z}) = -\frac{1}{2} \sum_{i,j} w_{ij} z_i z_j - \sum_i \theta_i z_i$$

Learning of autoassociative memory = placing of attractors on or near regions of interest

## Hopfield Learning Rule

$$w_{ij} = \frac{1}{N} \sum_{k \in data} z_i^k z_j^k$$

## Example

---

code/ann\_hopfieldnet.py

---

```
from npamlib import *

data = np.array([[1,1,-1,-1],[1,-1,1,-1]])

#"Train" Weight Matrix
W = np.zeros([len(data), 4, 4])
for i in range(len(data)):
    W[i, :, :] += np.outer(data[i], data[i].T)
W = np.mean(W, axis=0)
b = data.mean(axis=0)

out = FRN(N=4, b=b, W=W, rates_in = [1,1,1,-1], activation_function=hopfield, T=5)
```

---

## Energy

$$E(\mathbf{z}) = -\frac{1}{2} \sum_{i,j} w_{ij} z_i z_j - \sum_i \theta_i z_i$$

$z_0$	$z_1$	$z_2$	$z_3$	Energy
[-1	-1	-1	-1]	-0.0
[-1	-1	-1	1]	-2.0
[-1	-1	1	-1]	-4.0
[-1	-1	1	1]	-6.0
[-1	1	-1	-1]	-4.0
[-1	1	-1	1]	-6.0
[-1	1	1	-1]	-0.0
[-1	1	1	1]	-2.0
[ 1	-1	-1	-1]	-6.0
[ 1	-1	-1	1]	-0.0
[ 1	-1	1	-1]	-10.0
[ 1	-1	1	1]	-4.0
[ 1	1	-1	-1]	-10.0
[ 1	1	-1	1]	-4.0
[ 1	1	1	-1]	-6.0
[ 1	1	1	1]	-0.0

## Example Letters

code/ann\_hopfieldnet\_letter.py

---

```
from npamplib import *
data, labels = data_load_letters()
#Transform into +1/-1
data = (data>.5)*2-1

#"Train" Weight Matrix
W = np.zeros([len(data), 25, 25])
for i in range(len(data)):
    W[i, :, :] += np.outer(data[i], data[i].T)
W = np.mean(W, axis=0)-np.eye(25)
b = data.mean(axis=0)

res = np.zeros([len(data),25])
data_corrupt = np.empty_like(data)

#Number of pixels to corrupt
N_corrupt = 5
for i in range(len(data)):
    d = data[i,:]
    #Flip pixel value
    d[np.random.choice(range(25),N_corrupt)] *= -1
    data_corrupt[i,:] = d
    res[i] = FRN(N=25, b=b, W=W, rates_in = data[i], activation_function=hopfield, T=10)
```

---

## Example Letters

code/ann\_hopfieldnet\_letter.py

---

```
from npamplib import *
data, labels = data_load_letters()
#Transform into +1/-1
data = (data>.5)*2-1

#"Train" Weight Matrix
W = np.zeros([len(data), 25, 25])
for i in range(len(data)):
    W[i, :, :] += np.outer(data[i], data[i].T)
W = np.mean(W, axis=0)-np.eye(25)
b = data.mean(axis=0)

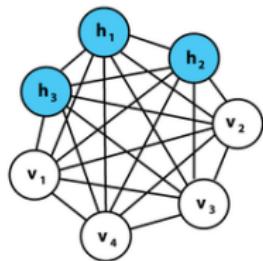
res = np.zeros([len(data),25])
data_corrupt = np.empty_like(data)

#Number of pixels to corrupt
N_corrupt = 5
for i in range(len(data)):
    d = data[i,:]
    #Flip pixel value
    d[np.random.choice(range(25),N_corrupt)] *= -1
    data_corrupt[i,:] = d
    res[i] = FRN(N=25, b=b, W=W, rates_in = data[i], activation_function=hopfield, T=10)
```

---

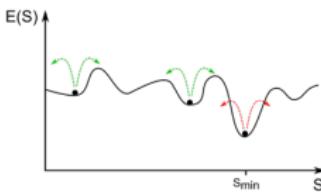
No labels were used for learning: Unsupervised Learning

## Stochastic Neural Networks - Boltzmann Machine (BM)

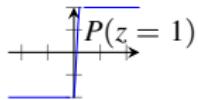


$$p^{BM}(\mathbf{z}) = \frac{1}{Z} \exp(-E(\mathbf{z})), \quad E(\mathbf{z}) = -\frac{1}{2} \sum_{i,j} w_{ij} z_i z_j - \sum_i \theta_i z_i$$

Intuition behind the BM: add stochasticity to a Hopfield network to escape local minima

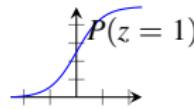


Deterministic unit (Hopfield Network)



$$\Delta w_{ij} \propto \langle z_i z_j \rangle_{data}$$

Stochastic unit (Boltzmann Machine)



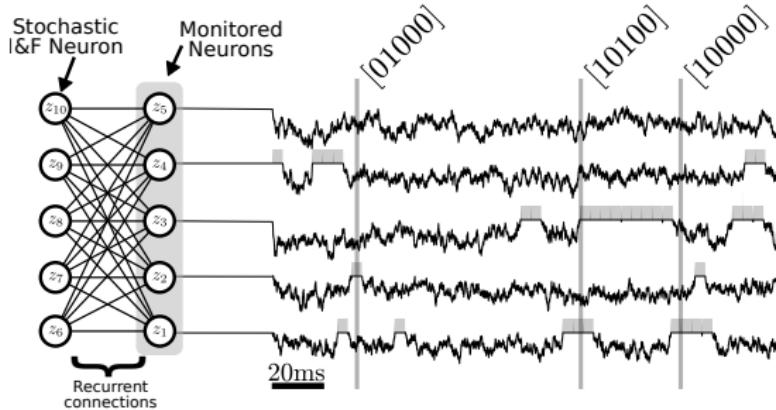
$$\Delta w_{ij} \propto \langle z_i z_j \rangle_{data} - \underbrace{\langle z_i z_j \rangle_{reconstruction}}_{\text{difficult to evaluate}}$$

Hinton and Sejnowski, 1986

The additional term is often estimated using Markov Chain Monte Carlo (MCMC) sampling

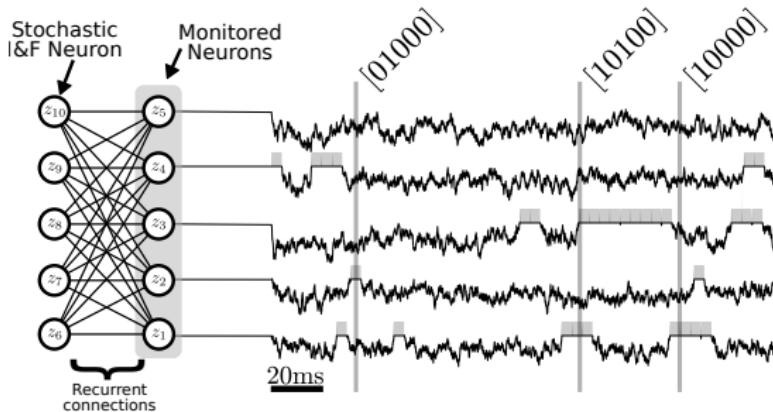
## Neural sampling and learning with noisy Integrate & Fire (I&F) neurons

Neuron state associated with a binary random variable  $z_k = 0, 1$

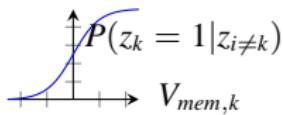


## Neural sampling and learning with noisy I&F neurons

Neuron state associated with a binary random variable  $z_k = 0, 1$



- Network states are Markov Chain Monte Carlo samples of the Boltzmann distribution if:



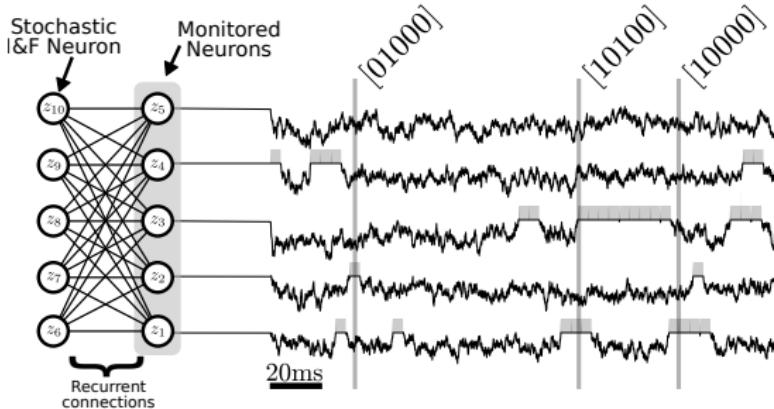
$$V_{mem,k}(t) = \sum_i W_{ki} z_i(t) + b_k$$

Linear sum of inputs

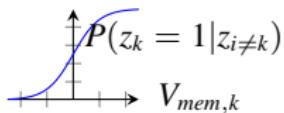
Buesing, Bill, Nessler, and Maass, *PLoS Computational Biology*, 2011

## Neural sampling and learning with noisy I&F neurons

Neuron state associated with a binary random variable  $z_k = 0, 1$



- Network states are Markov Chain Monte Carlo samples of the Boltzmann distribution if:



$$V_{mem,k}(t) = \sum_i W_{ki} z_i(t) + b_k$$

Linear sum of inputs

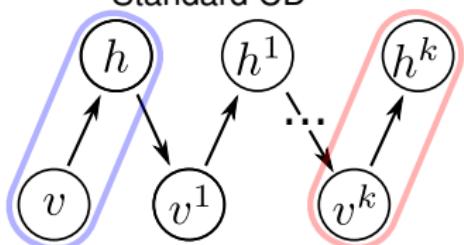
Buesing, Bill, Nessler, and Maass, *PLoS Computational Biology*, 2011

- Biophysical neuron models can meet these conditions
- Learning with STDP-based *Event-driven Contrastive Divergence*

Neftci, Posch, and Chicca, *Computational Intelligence*, 2014

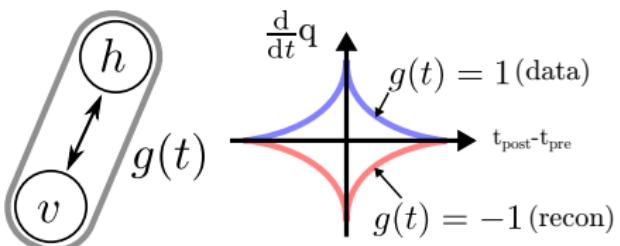
## Event-driven Contrastive Divergence

Gibbs Sampling +  
Standard CD



$$\Delta w \propto \langle vh \rangle_{\text{data}} - \langle v^k h^k \rangle_{\text{recon}}$$

Neural Sampling +  
Event-driven CD



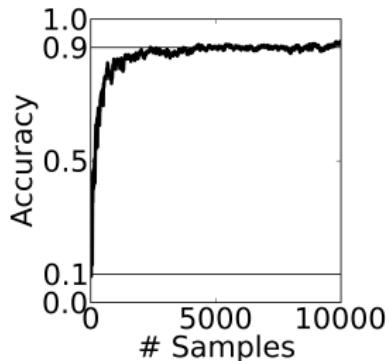
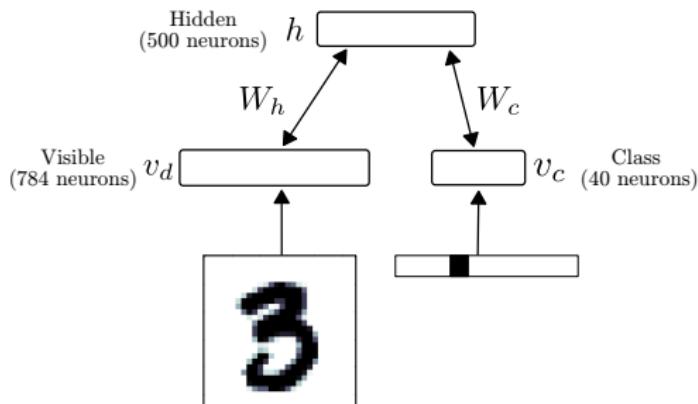
$$\frac{d}{dt}q = g(t) \text{STDP}(v(t), h(t))$$

In average, eCD behaves like standard CD

Neftci\_etal14 Neftci\_etal14 Neftci\_etal14 Neftci\_etal14

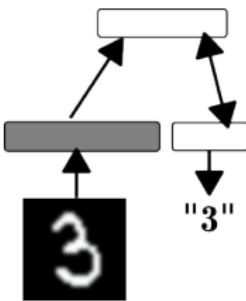
## Learning a model of MNIST hand-written digits (software simulations)

0 0 0 0 0 0 0 0 0 0 0 0 0  
1 1 1 1 1 1 1 1 1 1 1 1 1  
2 2 2 2 2 2 2 2 2 2 2 2 2  
3 3 3 3 3 3 3 3 3 3 3 3 3  
4 4 4 4 4 4 4 4 4 4 4 4 4  
5 5 5 5 5 5 5 5 5 5 5 5 5  
6 6 6 6 6 6 6 6 6 6 6 6 6  
7 7 7 7 7 7 7 7 7 7 7 7 7  
8 8 8 8 8 8 8 8 8 8 8 8 8  
9 9 9 9 9 9 9 9 9 9 9 9 9

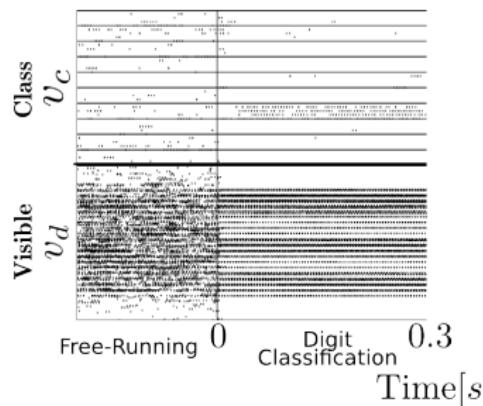


Recognition Accuracy: 91.9%, (vs. standard CD: 93.6%)

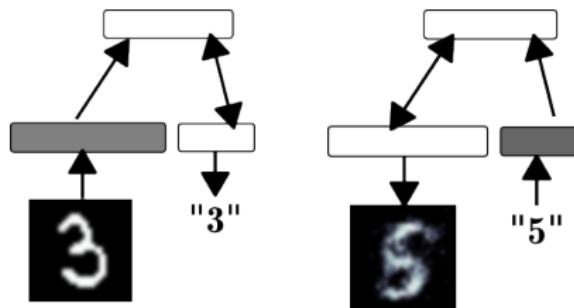
# The trained network performs Bayesian inference



Digit  
Classification

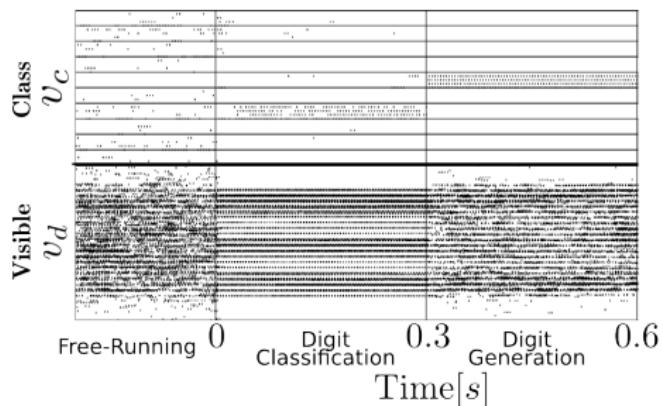


# The trained network performs Bayesian inference

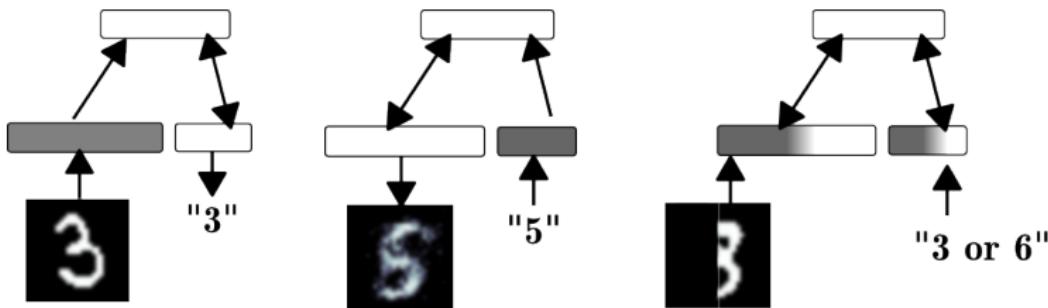


Digit  
Classification

Digit  
Generation



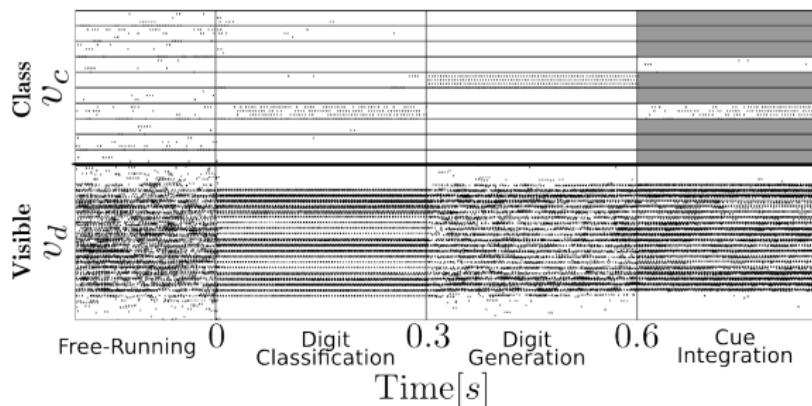
# The trained network performs Bayesian inference



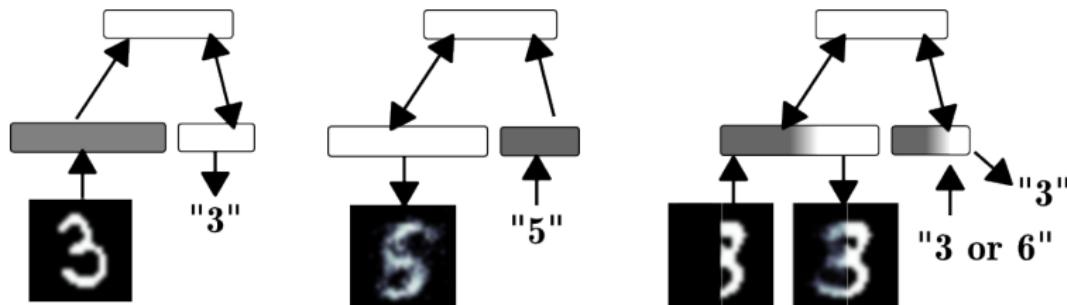
Digit  
Classification

Digit  
Generation

Cue  
Integration



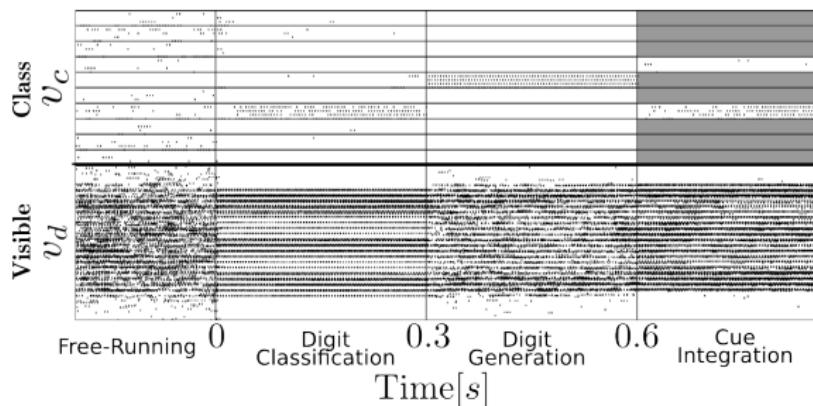
# The trained network performs Bayesian inference



Digit  
Classification

Digit  
Generation

Cue  
Integration

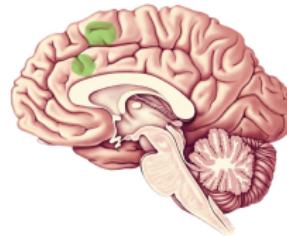
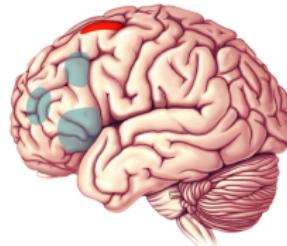


## Working Memory

# Working Memory

## Working Memory:

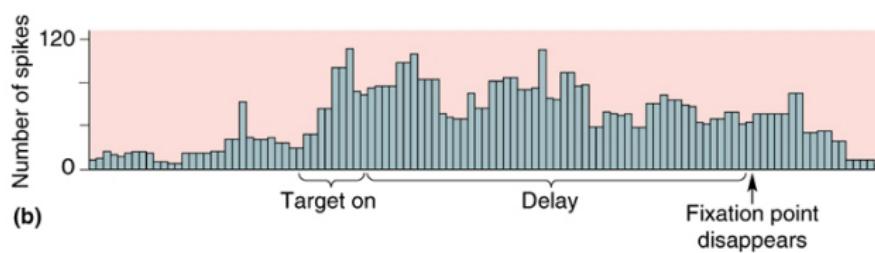
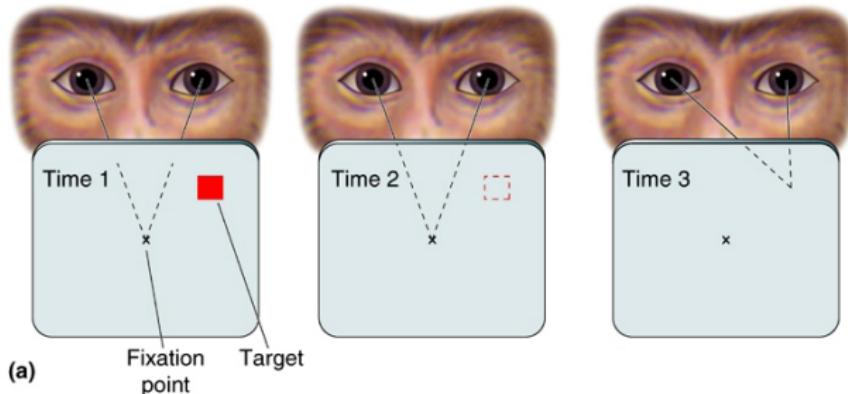
- A type of short-term memory
- Limited in capacity
- Task- and sensory modality-dependent
- Necessary for cognitive control



*Human brain areas for working memory of face identity and location*

Working Memory is a "Mental Sketch Pad"

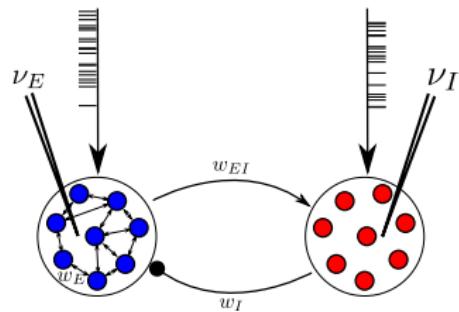
# Neural Correlates of Working Memory



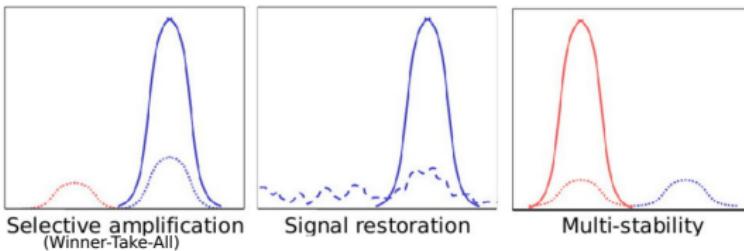
Neuroscience: Exploring the Brain, 3rd Ed, Bear, Connors, and Paradiso Copyright © 2007 Lippincott Williams & Wilkins

# Neural Models of Working Memory

## Excitatory Inhibitory Network



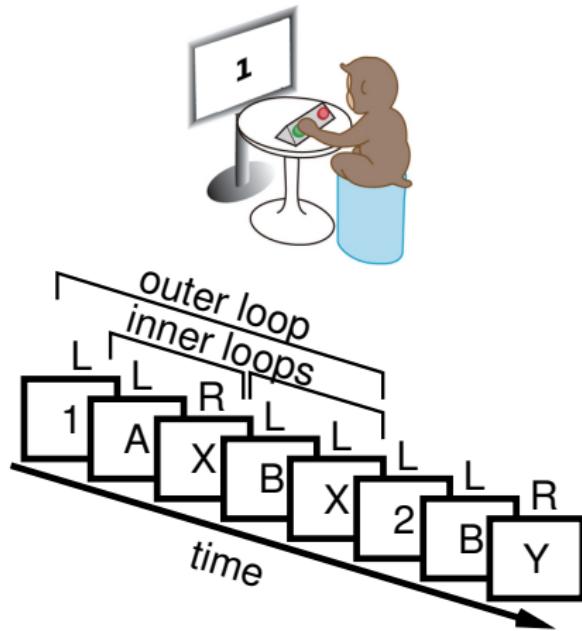
## Winner-Take-All (WTA) Properties



Douglas and Martin, *Annual Review of Neuroscience*, 2004

# Working Memory for Cognitive Control

The 12-AX Cognitive task



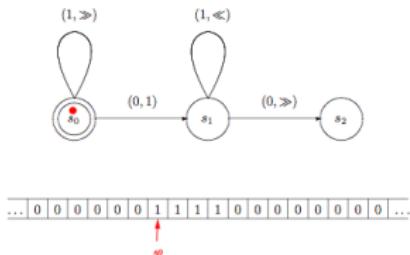
O'Reilly and Frank, 2006

## **Finite State Machines (FSM)**

## State Machines are the Main Components of Turing Machines

*A Turing machine is a theoretical computing machine invented by Alan Turing (1937) to serve as an idealized model for mathematical calculation.*

Wolfram Mathworld

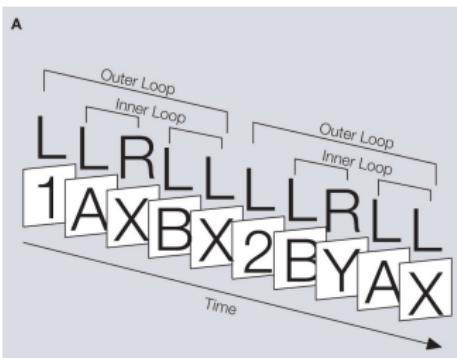


<http://www.google.com/doodles/alan-turings-100th-birthday>

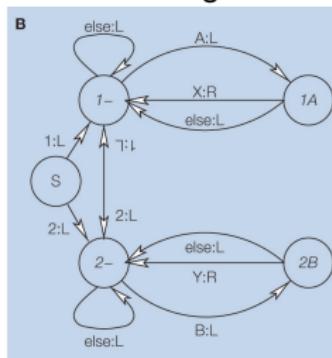
Any real-world computation can be translated into an equivalent computation involving a Turing machine

# The 12-AX Cognitive task

## 12AX Task



## State machine solving 12AX

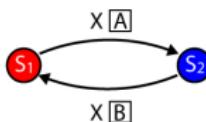


O'Reilly and Frank, Dayan, 2006, 2008

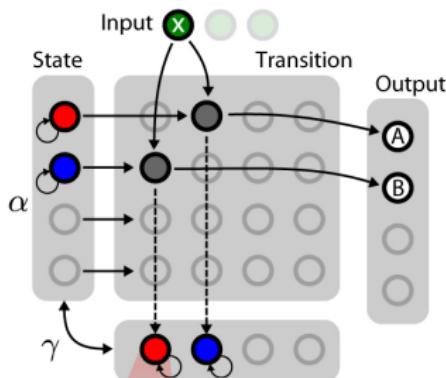
Can neurons implement state machines?

# Synthesizing State-Dependent Behavior in WTA

**(a) High-Level Behavioral Model**



**(b) Abstract Computational Layer**

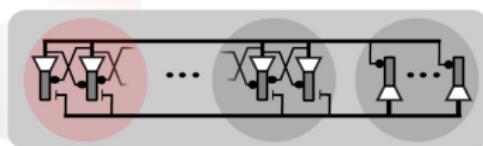


**(c) Neuronal Hardware**

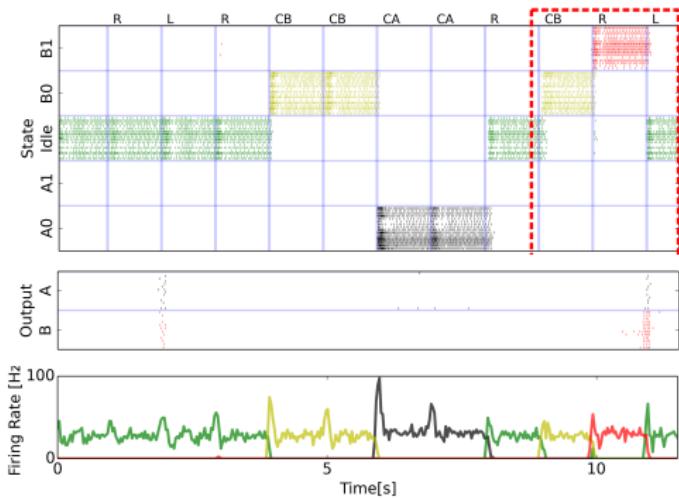
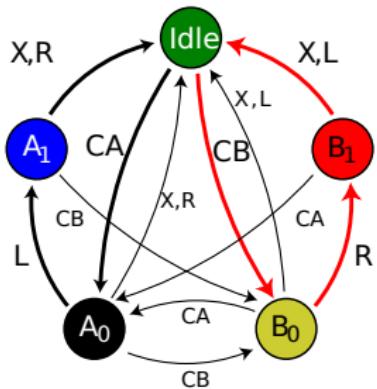
- Neuron
- Exc. Synapse
- Inh. Synapse



Inhibitory Population

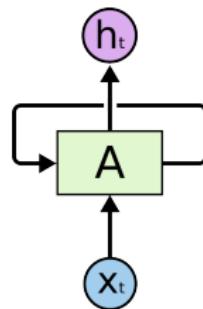


## Example: Synthesis of 12-AX task state machine

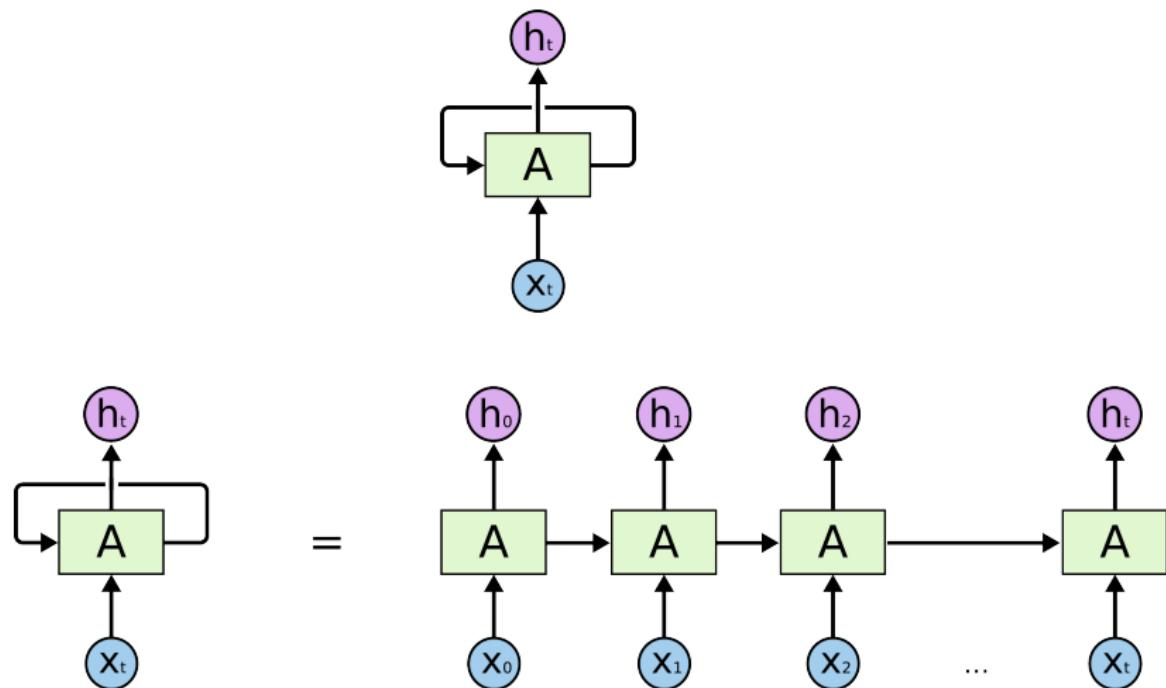


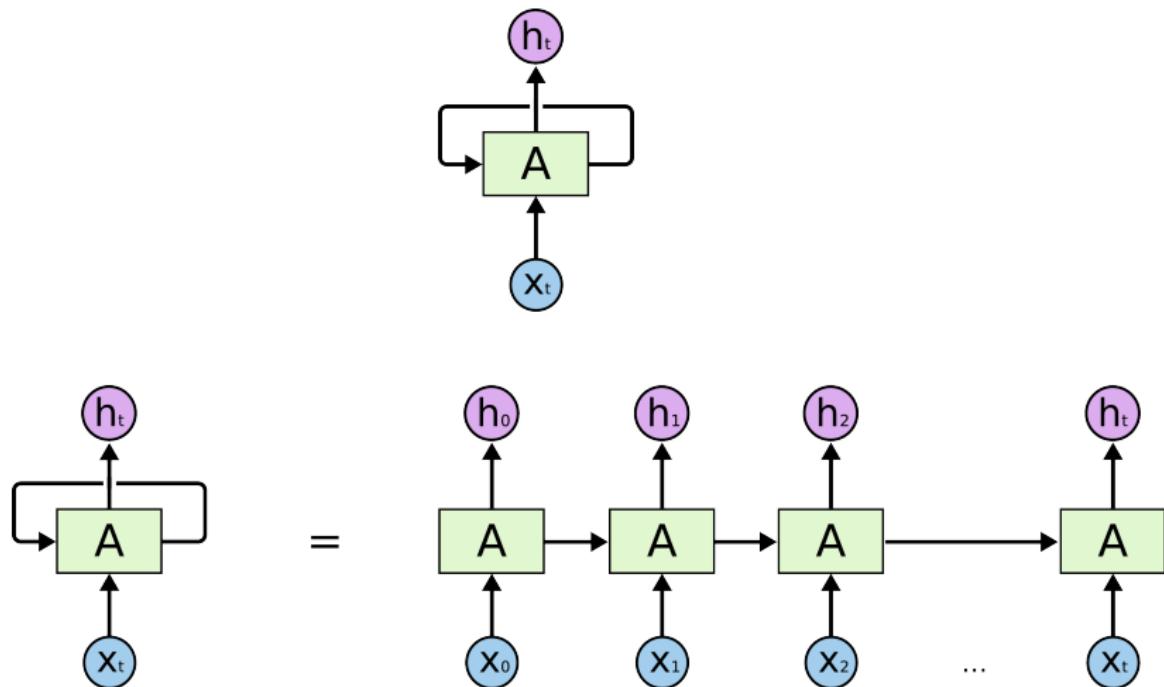
Neftci, Binas, Rutishauser, Chicca, Indiveri, and Douglas, PNAS, 2013

## Recurrent Neural Networks



## Recurrent Neural Networks





Learning with Back-Propagation on the “unrolled” network

- Topic in brain-inspired machine learning using what you have learned in class
- No class Nov. 25
- Dec 2: 10+5 min presentation
- Dec 7: Project report, using NIPS 2015 template <https://nips.cc/Conferences/2015/PaperInformation/StyleFiles>
- Discuss topic with instructor
- Groups of two OK, but delineate contributions